

# Big Data Management | Irish M50 Traffic Sensors Data

## Extract Transform Load | Apache Spark | Cassandra

This document explores ETL process using Apache Spark Data Frame API to perform data transformation queries and store the results in Cassandra Data Structures. Irish Road Network has installed sensors on specific location that collect vehicle data including location, type, speed, weight and more. This data is stored in atomic form.

**Dataset:** [Irish M50 Traffic Sensors Data](#)

Create a new keyspace named sensor.

```
cqlsh> describe keyspaces;

collections      streaming      system_schema  traffic
griffithcollege system         system_traces  training
"GriffithCollege" system_auth    system_views
stream          system_distributed system_virtual_schema

cqlsh> create keyspace sensors with replication =
... { 'class' : 'SimpleStrategy', 'replication_factor' : 1 };
cqlsh> describe keyspaces;

collections      stream      system_distributed system_virtual_schema
griffithcollege streaming   system_schema      traffic
"GriffithCollege" system     system_traces      training
sensors         system_auth system_views

cqlsh> use sensors;
```

Load dataset.

```
In [1]: import pyspark.sql.functions as f
        from pyspark.sql.functions import col

In [2]: # Create path for data file and load data
        path = "traffic/per-vehicle-records-2021-01-15.csv"
        data = spark.read.format("csv").options(header=True).load(path)

In [3]: # Get names of columns for future reference
        print(data.columns)

['cosit', 'year', 'month', 'day', 'hour', 'minute', 'second', 'millisecond', 'minuteofday', 'lane', 'lanename', 'straddlelane', 'straddlelanename', 'class', 'classname', 'length', 'headway', 'gap', 'speed', 'weight', 'temperature', 'duration', 'validitycode', 'numberofaxles', 'axleweights', 'axlespacings']
```

Create data structure in Cassandra.

```
cqlsh:sensors> CREATE TABLE percentage_per_vehicle_category (
... id int primary key,
... category text,
... percentage float
... );_
```

## Perform Computation in Pyspark.

```
In [4]: # 1. Calculate the usage of Irish road network in terms of percentage grouped
# by vehicle category.
```

```
total = data.count()
q1 = data.groupby('classname').count().withColumnRenamed('classname', 'category')\
.withColumn('percentage', f.round(f.col('count')/total*100, 2)).drop('count')\
.withColumn('id', f.monotonically_increasing_id())
q1.show()
```

category	percentage	id
CAR	70.24	0
HGV_ART	7.57	1
BUS	0.78	2
HGV_RIG	4.37	3
null	0.01	4
CARAVAN	0.62	5
LGV	15.84	6
MBIKE	0.56	7

## Store dataframe to Cassandra and Check.

### Move above data frame to Cassandra

```
In [6]: # Move above data frame to Cassandra
q1.select("id", "category", "percentage")\
.write.format("org.apache.spark.sql.cassandra")\
.options(table="percentage_per_vehicle_category", keyspace="sensors")\
.save(mode="append")
```

### Read stored data frame from Cassandra

```
In [7]: # Read stored data frame from Cassandra
spark.read.format("org.apache.spark.sql.cassandra")\
.load(keyspace='traffic', table='percentage_per_vehicle_category').orderBy('id').show()
```

id	category	percentage
0	CAR	70.24
1	HGV_ART	7.57
2	BUS	0.78
3	HGV_RIG	4.37
4	null	0.01
5	CARAVAN	0.62
6	LGV	15.84
7	MBIKE	0.56

## Checking in Cassandra.

```
cqlsh:sensors> SELECT * FROM percentage_per_vehicle_category;
```

id	category	percentage
5	CARAVAN	0.62
1	HGV_ART	7.57
0	CAR	70.24
2	BUS	0.78
4	null	0.01
7	MBIKE	0.56
6	LGV	15.84
3	HGV_RIG	4.37

(8 rows)

## Create location dictionaries with sensor codes.

These are locations for each sensor of motorway junctions.

```
In [8]: # Get cosit for each junction for motorways
# Note: Cosit for Jn01-Jn02 Dublin port to Santry is not available on the site map
m50 = {'000000001012': 'Jn02-Jn03',
       '000000001500': 'Jn03-Jn04',
       '000000001501': 'Jn04-Jn05',
       '000000001502': 'Jn05-Jn06',
       '000000001508': 'Jn06-Jn07',
       '000000001503': 'Jn07-Jn09',
       '000000001509': 'Jn09-Jn10',
       '000000001504': 'Jn10-Jn11',
       '000000001505': 'Jn11-Jn12',
       '000000001506': 'Jn12-Jn13',
       '000000001507': 'Jn13-Jn14',
       '0000000015010': 'Jn14-Jn15',
       '0000000015011': 'Jn15-Jn16',
       '0000000015012': 'Jn16-Jn17'
      }
```

Create cassndra structure for motorway hourly flows.

```
cqlsh:sensors> CREATE TABLE motorway_hourly_flows (
... flow text,
... hour int,
... vehicle_count int,
... PRIMARY KEY (flow)
... );
```

Perform computation.

## # 2. Calculate the highest and lowest hourly flows on M50 - show the

```
In [9]: # 2. Calculate the highest and lowest hourly flows on M50 - show the
# hours and total number of vehicle counts

hourly_flows = data.select("cosit", "hour").where(f.col('cosit').isin(list(m50.keys()))).groupBy('hour').count().sort('count')
mx = hourly_flows.agg({'count' : 'max'}).collect()[0][0] # Collect Max Value
mn = hourly_flows.agg({'count' : 'min'}).collect()[0][0] # Collect Min value

# Select row where max value
mx_flow = hourly_flows.select('hour', 'count').where(f.col('count') == mx).withColumn('flow', f.lit('highest'))
# Select row where min value
mn_flow = hourly_flows.select('hour', 'count').where(f.col('count') == mn).withColumn('flow', f.lit('lowest'))

# Union both max and min rows to form a table
motorway_hourly_flows = mx_flow.union(mn_flow).withColumnRenamed('count', 'vehicle_count')
motorway_hourly_flows.show() # add auto_increment column for ids
```

hour	vehicle_count	flow
16	38655	highest
2	1167	lowest

Move dataframe to Cassandra and check.

## Move above data frame to Cassandra

```
In [11]: motorway_hourly_flows.select("flow", "hour", "vehicle_count")\
.write.format("org.apache.spark.sql.cassandra")\
.options(table="motorway_hourly_flows", keyspace="sensors")\
.save(mode="append")

In [12]: spark.read.format("org.apache.spark.sql.cassandra")\
.load(keyspace='traffic', table='motorway_hourly_flows').show()
```

flow	hour	vehicle_count
lowest	2	1167
highest	16	38655

Check in Cassandra.

```
cqlsh:sensors> SELECT * FROM motorway_hourly_flows;

flow      | hour | vehicle_count
-----+-----+-----
lowest    | 2    | 1167
highest   | 16   | 38655

(2 rows)
cqlsh:sensors> _
```

Create Cassandra data structure for motorway rush hours.

```
cqlsh:sensors> CREATE TABLE motorway_rush_hours (
    ... flow text,
    ... hour int,
    ... vehicle_count int,
    ... PRIMARY KEY (flow)
    ... );
```

Perform Computation.

### 3. Calculate the evening and morning rush hours on M50

```
In [8]: # 3. Calculate the evening and morning rush hours on M50 - show the
# hours and the total counts.

morning_hours = ['6','7','8','9','10']
evening_hours = ['16','17','18','19','20']

data.head()
morning_flows = data.select('hour', 'cosit').where(f.col('cosit').isin(list(m50.keys())))\
.where(f.col('hour').isin(morning_hours))\
.groupby('hour').count().withColumn('time', f.lit('morning'))

evening_flows = data.select('hour', 'cosit').where(f.col('cosit').isin(list(m50.keys())))\
.where(f.col('hour').isin(evening_hours))\
.groupby('hour').count().withColumn('time', f.lit('evening'))

rush_hours = morning_flows.union(evening_flows)

rush_hours.show()
```

```
+---+-----+-----+
|hour|count|  time|
+---+-----+-----+
| 7|22528|morning|
| 8|27180|morning|
| 6|18728|morning|
| 9|29992|morning|
|10|29279|morning|
|16|38655|evening|
|18|18173|evening|
|17|36016|evening|
|19|13788|evening|
|20|11647|evening|
+---+-----+-----+
```

Move Dataframe to Cassandra and Check.

```
In [9]: rush_hours.select("hour", "count", "time")\
        .write.format("org.apache.spark.sql.cassandra")\
        .options(table="motorway_rush_hour", keyspace="sensors")\
        .save(mode="append")

In [11]: spark.read.format("org.apache.spark.sql.cassandra")\
        .load(keyspace='sensors', table='motorway_rush_hour').show()

+-----+-----+-----+
|hour|count|  time|
+-----+-----+-----+
| 18|18173|evening|
|  9|29992|morning|
| 17|36016|evening|
| 20|11647|evening|
|  7|22528|morning|
| 10|29279|morning|
| 16|38655|evening|
| 19|13788|evening|
|  8|27180|morning|
|  6|18728|morning|
+-----+-----+-----+
```

Check in Cassandra.

```
ali@bdm:~$ cqlsh
Connected to Test Cluster at 127.0.0.1:9042
[cqlsh 6.1.0 | Cassandra 4.1 | CQL spec 3.4.6 | Native protocol v5]
Use HELP for help.
cqlsh> use sensors;
cqlsh:sensors> SELECT * FROM motorway_rush_hour;

hour | count |  time
-----+-----+-----
 10 | 29279 | morning
 16 | 38655 | evening
 19 | 13788 | evening
  8 | 27180 | morning
 18 | 18173 | evening
 20 | 11647 | evening
  7 | 22528 | morning
  6 | 18728 | morning
  9 | 29992 | morning
 17 | 36016 | evening

(10 rows)
cqlsh:sensors> _
```

Create a Cassandra data structure for motorway average speed.

```
cqlsh:sensors> CREATE TABLE motorway_average_speed ( cosit int, location text, average_speed float,
PRIMARY KEY (cosit) );_
```

Perform Computation.

First get average speed by cosit.

#### # 4. Calculate average speed between each junction on M50

```
In [12]: # 4. Calculate average speed between each junction on M50 (e.g., junction
# 1 - junction2, junction 2 - junction 3, etc.).

motorway_data = data.where(f.col('cosit').isin(list(m50.keys()))))
average_speed = motorway_data.select('cosit', col('speed').cast('double')).alias('speed')).groupby('cosit').mean('speed')
average_speed = average_speed.withColumnRenamed('avg(speed)', 'average')
average_speed.show()
```

cosit	average
000000001500	88.83526554404145
000000001501	104.02299711199059
000000001505	98.92545893412945
000000001503	98.45699912510936
000000001509	94.73736586836881
000000001502	99.01588546773877
000000001507	102.64251095162643
000000001506	102.11667798306114
000000001501	98.10988853617204
000000001012	84.09989342515166
0000000015010	106.05619648259243
0000000015012	106.45533712709087
000000001504	100.41781593019984
000000001508	96.13615310118321

Than get location by cosit.

```
In [13]: average_speed.createTempView('average_speed')
location = spark.createDataFrame(data=m50.items(), schema=['cosit', 'location'])
location.createTempView('location')
location.show()
```

cosit	location
000000001012	Jn02-Jn03
000000001500	Jn03-Jn04
000000001501	Jn04-Jn05
000000001502	Jn05-Jn06
000000001508	Jn06-Jn07
000000001503	Jn07-Jn09
000000001509	Jn09-Jn10
000000001504	Jn10-Jn11
000000001505	Jn11-Jn12
000000001506	Jn12-Jn13
000000001507	Jn13-Jn14
0000000015010	Jn14-Jn15
0000000015011	Jn15-Jn16
0000000015012	Jn16-Jn17

Then joining both tables with cosit as key.

```
In [15]: motorway_average_speed = spark.sql("SELECT average_speed.cosit, location, ROUND(average, 2) as average_speed \
FROM average_speed, location\
WHERE average_speed.cosit = location.cosit ORDER BY location")
motorway_average_speed.show()
```

cosit	location	average_speed
000000001012	Jn02-Jn03	84.1
000000001500	Jn03-Jn04	88.84
000000001501	Jn04-Jn05	98.11
000000001502	Jn05-Jn06	99.02
000000001503	Jn06-Jn07	96.14
000000001503	Jn07-Jn09	98.46
000000001509	Jn09-Jn10	94.74
000000001504	Jn10-Jn11	100.42
000000001505	Jn11-Jn12	98.93
000000001506	Jn12-Jn13	102.12
000000001507	Jn13-Jn14	102.64
0000000015010	Jn14-Jn15	106.06
0000000015011	Jn15-Jn16	104.02
0000000015012	Jn16-Jn17	106.46

Move to Cassandra and Check.

```
In [16]: motorway_average_speed.select("cosit", "location", "average_speed")\
        .write.format("org.apache.spark.sql.cassandra")\
        .options(table="motorway_average_speed", keyspace="sensors")\
        .save(mode="append")
```

```
In [17]: spark.read.format("org.apache.spark.sql.cassandra")\
        .load(keyspace="sensors", table='motorway_average_speed').show()
```

cosit	average_speed	location
1507	102.64	Jn13-Jn14
1508	96.14	Jn06-Jn07
1505	98.93	Jn11-Jn12
15011	104.02	Jn15-Jn16
1504	100.42	Jn10-Jn11
1501	98.11	Jn04-Jn05
15010	106.06	Jn14-Jn15
1509	94.74	Jn09-Jn10
1012	84.1	Jn02-Jn03
1502	99.02	Jn05-Jn06
1506	102.12	Jn12-Jn13
15012	106.46	Jn16-Jn17
1503	98.46	Jn07-Jn09
1500	88.84	Jn03-Jn04

Check in Cassandra.

```
cqlsh:sensors> SELECT * FROM motorway_average_speed;
```

cosit	average_speed	location
1505	98.93	Jn11-Jn12
1500	88.84	Jn03-Jn04
15011	104.02	Jn15-Jn16
1504	100.42	Jn10-Jn11
1506	102.12	Jn12-Jn13
15012	106.46	Jn16-Jn17
1503	98.46	Jn07-Jn09
1012	84.1	Jn02-Jn03
1501	98.11	Jn04-Jn05
15010	106.06	Jn14-Jn15
1509	94.74	Jn09-Jn10
1502	99.02	Jn05-Jn06
1507	102.64	Jn13-Jn14
1508	96.14	Jn06-Jn07

```
(14 rows)
cqlsh:sensors>
```